

Problem xy (use separate files for each problem)

10111

13 januar, 2023

```
install.packages("knitr")
install.packages("MASS")
install.packages("caret")
install.packages("pls")
install.packages("glmnet")
install.packages("gam")
install.packages("gbm")
install.packages("randomForest")
install.packages("ggfortify")
install.packages("leaps")
install.packages("pROC")
install.packages("sfsmisc")
```

```
id <- "1kGOLsnKA0Uq2lWKlMjhAF8h71sc0WcLO" # google file ID
d.bodyfat <- read.csv(sprintf("https://docs.google.com/uc?id=%s&export=download",
  id))[, -c(1)]
set.seed(1234)
training_set_size <- floor(0.8 * nrow(d.bodyfat))

samples <- sample(1:nrow(d.bodyfat), training_set_size, replace = F)
d.body.train <- d.bodyfat[samples, ]
d.body.test <- d.bodyfat[-samples, ]
```

a)

```
r.lm.BodyFat <- lm(BodyFat ~ . - Abdomen + poly(Abdomen, degree = 2), d.body.train)
summary(r.lm.BodyFat)
```

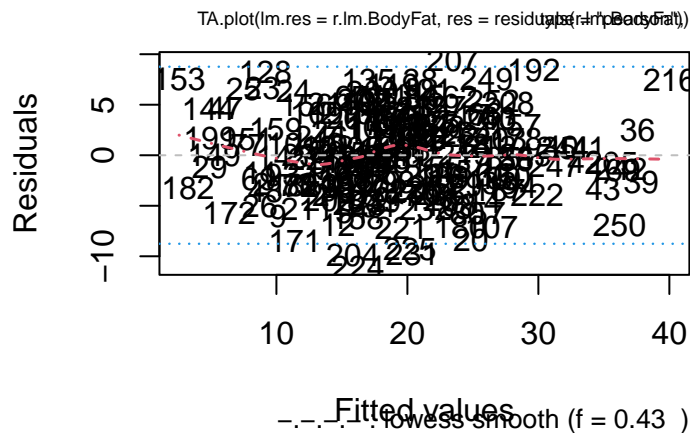
```
##
## Call:
## lm(formula = BodyFat ~ . - Abdomen + poly(Abdomen, degree = 2),
##     data = d.body.train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -11.0198  -2.9100  -0.1409   2.9595   9.3920
##
## Coefficients:
```

```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)    80.91992    21.31906   3.796 0.000199 ***
## Age            0.06754     0.03566   1.894 0.059765 .
## Weight        -0.04196     0.06164  -0.681 0.496910
## Height        -0.06602     0.10086  -0.655 0.513560
## Neck          -0.55179     0.25729  -2.145 0.033285 *
## Chest         -0.08479     0.10716  -0.791 0.429830
## Hip           -0.09304     0.16728  -0.556 0.578760
## Thigh          0.08789     0.15638   0.562 0.574768
## Knee          -0.10641     0.27650  -0.385 0.700799
## Ankle          0.11223     0.23063   0.487 0.627087
## Biceps         0.35094     0.20161   1.741 0.083391 .
## Forearm        0.33282     0.21528   1.546 0.123807
## Wrist         -2.08478     0.58018  -3.593 0.000418 ***
## poly(Abdomen, degree = 2)1 138.18455    15.33920   9.009 2.49e-16 ***
## poly(Abdomen, degree = 2)2 -12.58157     5.04401  -2.494 0.013490 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.234 on 186 degrees of freedom
## Multiple R-squared:  0.7559, Adjusted R-squared:  0.7375
## F-statistic: 41.15 on 14 and 186 DF, p-value: < 2.2e-16
```

The R^2 is 0.7559. This is rather large, but we see that the Adjusted R^2 is smaller. This may indicate that the additional variables in the model is not adding value to the model.


```
TA.plot(r.lm.BodyFat, res = residuals(r.lm.BodyFat, type = "pearson"))
```

BodyFat ~ . - Abdomen + poly(Abdomen, c



```
qqnorm(d.body.train$BodyFat, pch = 1, frame = FALSE)
qqline(d.body.train$BodyFat, col = "steelblue", lwd = 2)
```

Normal Q-Q Plot



```
# qqplot(d.body.train$BodyFat)
```

b)


Below you have to complete the code and then replace `eval=FALSE` by `eval=TRUE` in the chunk options:

```
regfit_fwd = regsubsets(BodyFat ~ . - Abdomen + poly(Abdomen, degree = 2), data = d.body.train,
  nvmax = 14, method = "forward")
regfit_fwd_summary <- summary(regfit_fwd)
regfit_fwd_summary$outmat
```

##		Age	Weight	Height	Neck	Chest	Hip	Thigh	Knee	Ankle	Biceps	Forearm
## 1	(1)	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
## 2	(1)	" "	"*	" "	" "	" "	" "	" "	" "	" "	" "	" "
## 3	(1)	" "	"*	" "	" "	" "	" "	" "	" "	" "	" "	" "
## 4	(1)	" "	"*	" "	" "	" "	" "	" "	" "	" "	" "	" "
## 5	(1)	" "	"*	" "	" "	" "	" "	" "	" "	" "	"*	" "
## 6	(1)	" "	"*	" "	"*	" "	" "	" "	" "	" "	"*	" "
## 7	(1)	"*	"*	" "	"*	" "	" "	" "	" "	" "	"*	" "
## 8	(1)	"*	"*	" "	"*	" "	" "	" "	" "	" "	"*	"*
## 9	(1)	"*	"*	" "	"*	"*	" "	" "	" "	" "	"*	"*
## 10	(1)	"*	"*	"*	"*	"*	" "	" "	" "	" "	"*	"*
## 11	(1)	"*	"*	"*	"*	"*	"*	" "	" "	" "	"*	"*
## 12	(1)	"*	"*	"*	"*	"*	"*	"*	" "	" "	"*	"*
## 13	(1)	"*	"*	"*	"*	"*	"*	"*	" "	"*	"*	"*
## 14	(1)	"*	"*	"*	"*	"*	"*	"*	"*	"*	"*	"*
##		Wrist poly(Abdomen, degree = 2)1 poly(Abdomen, degree = 2)2										
## 1	(1)	" "	"*					" "				
## 2	(1)	" "	"*					" "				
## 3	(1)	"*	"*					" "				
## 4	(1)	"*	"*					"*				
## 5	(1)	"*	"*					"*				
## 6	(1)	"*	"*					"*				

```
## 7 ( 1 ) "*" "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
## 11 ( 1 ) "*" "*"
## 12 ( 1 ) "*" "*"
## 13 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "*"
## 14 ( 1 ) "*" "
```

```
plot(regfit_fwd_summary$bic, main = "Forward Stepwise Selection", xlab = "Number of Predictors",
     ylab = "BIC", type = "l")
```



Here we can see that the model with 4 predictors give the lowest BIC. That is model 6 with predictors Weight, Neck, Biceps and Wrist.

```
best_model <- lm(BodyFat ~ Weight + Neck + Biceps + Wrist, d.body.train)

mse.best_model = mean((d.body.test$BodyFat - predict(best_model, newdata = d.body.test))^2)
mse.best_model
```


```
## [1] 40.89131
```

The MSE is 40.89131 with the reduced model

c)

```
x.train <- model.matrix(BodyFat ~ . - Abdomen + poly(Abdomen, degree = 2), data = d.body.train)[,
-1]
y.train <- d.body.train$BodyFat
x.test = model.matrix(BodyFat ~ . - Abdomen + poly(Abdomen, degree = 2), data = d.body.test)[,
-1]
y.test = d.body.test$BodyFat
```

```
set.seed(4268)
cv.lasso <- cv.glmnet(x.train, y.train, alpha = 1)
plot(cv.lasso)
```



```
cv.lasso$lambda.1se
```

```
## [1] 0.4455843
```

```
bodyfat.lasso <- glmnet(x.train, y.train, alpha = 1, lambda = cv.lasso$lambda.1se)
coef(bodyfat.lasso)
```

```
## 15 x 1 sparse Matrix of class "dgCMatrix"
##                               s0
## (Intercept)                47.09872944
## Age                        0.04167346
## Weight                      .
## Height                   -0.12282724
## Neck                      .
## Chest                     .
## Hip                      .
## Thigh                    .
## Knee                     .
## Ankle                    .
## Biceps                   .
## Forearm                  .
## Wrist                   -1.14587023
## poly(Abdomen, degree = 2)1  95.63871126
## poly(Abdomen, degree = 2)2 -11.14529896
```

```
mse.lasso = mean((y.test - predict(bodyfat.lasso, newx = x.test))^2)
mse.lasso
```

```
## [1] 50.31623
```

The MSE for Lasso is 50.31623.

```
cv.lasso$lambda.min
```

```
## [1] 0.09163496
```

```
bodyfat.lasso.min <- glmnet(x.train, y.train, alpha = 1, lambda = cv.lasso$lambda.min)
coef(bodyfat.lasso.min)
```

```
## 15 x 1 sparse Matrix of class "dgCMatrix"
##                                     s0
## (Intercept)                      69.29797647
## Age                               0.07085015
## Weight                             .
## Height                           -0.11547703
## Neck                             -0.42359496
## Chest                             .
## Hip                               .
## Thigh                             .
## Knee                             -0.08296349
## Ankle                             .
## Biceps                           0.17064055
## Forearm                          0.20158403
## Wrist                            -2.02171921
## poly(Abdomen, degree = 2)1 112.32782320
## poly(Abdomen, degree = 2)2 -14.66393600
```

```
mse.lasso.min = mean((y.test - predict(bodyfat.lasso.min, newx = x.test))^2)
mse.lasso.min
```

```
## [1] 74.07258
```

The MSE when λ_{min} is used is much larger than when λ_{1se} is used, thus the model with λ_{min} is not useful.

d)

```
# Had problems running this in markdown but it worked in rscript
pca.train <-
# prcomp(d.body.train[, c(2:6,8:-1)] + poly(d.body.train[, c(7)], degree = 2),
# scale = TRUE) var_explained = pca.train$sdev^2 / sum(pca.train$sdev^2)

# screeplot(pca.train, npcs = min(13, length(pca.train$sdev)), type =
# c('lines'))
```

See that the first PC explain over 60% of variability, and PC 2-4 explain approximately 10%, and the rest go near 0. So a useful number would be 5 PCs.

```
pcr_fit = pcr(BodyFat ~ . - Abdomen + poly(Abdomen, degree = 2), data = d.body.train,
  scale = TRUE, validation = "CV")
summary(pcr_fit)
```

```
## Data:      X dimension: 201 14
## Y dimension: 201 1
## Fit method: svdpc
## Number of components considered: 14
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV              8.285   6.671   5.698   5.632   5.538   5.324   5.704
## adjCV           8.285   6.665   5.656   5.628   5.519   5.371   5.667
##      7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV          5.509   5.102   4.954   4.914   4.875   5.021   4.558
## adjCV       5.463   5.054   4.926   4.889   4.851   4.985   4.525
##      14 comps
## CV          4.476
## adjCV       4.451
##
## TRAINING: % variance explained
##      1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8 comps
## X          56.69   66.41   75.17   81.79   86.37   90.61   92.83   94.91
## BodyFat     36.12   54.98   55.53   58.39   58.40   62.92   66.50   69.40
##      9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X          96.67   98.00   98.96   99.50   99.84   100.00
## BodyFat     69.87   70.18   70.70   71.38   75.53   75.59
```

```
validationplot(pcr_fit, val.type = "MSEP")
```

