

TMA4300 Computer Intensive Statistical Methods

Exercise 3, Spring 2023

Note: The report should consist of one (and only one) pdf-file, and should be uploaded via Blackboard, as specified in the course home page, no later than **Sunday, April 30th, 11:59 pm**. The report should include derivation of formulas that you are using in your implementations, the R code you have used to solve the project and the plots you have generated. Associated to the various plots there should be captions explaining the contents of the plots, and in addition all the plots should be explained and discussed in the main text of the report.

The project report should be formulated as a scientific report. In particular, it should be possible to understand what you have done without reading the questions in this problem text. Moreover, the text in the project report should consist of full sentences and proper punctuation should be used throughout. All results you present should also be discussed, what can you (and the world) learn from your results?

Getting started: The aim of this exercise is to get experience with the techniques we learned in part 3 of the course. As always make sure that each part of the code runs properly and discuss your findings.

Hint: In the bootstrap exercises you will need to use the R-function `sample`.

Problem A: Comparing $AR(2)$ parameter estimators using resampling of residuals

The data files and pre-programmed R-code can be downloaded from the course webpage. Look in the `probAhelp.R`-file and read the documentation to see how the code works. Load the code and data into R with

```
source("probAhelp.R")
source("probAdata.R")
```

In this exercise you should analyse the data in `data3A$x`, which contains a sequence of length $T = 100$ of a non-Gaussian time-series, and compare two different parameter estimators.

We consider an $AR(2)$ model which is specified by the relation

$$x_t = \beta_1 x_{t-1} + \beta_2 x_{t-2} + e_t,$$

where e_t are iid random variable with zero mean and constant variance.

The least sum of squared residuals (LS) and least sum of absolute residuals (LA) are obtained by minimising the following loss functions with respect to β :

$$Q_{LS}(\mathbf{x}) = \sum_{t=3}^T (x_t - \beta_1 x_{t-1} - \beta_2 x_{t-2})^2$$
$$Q_{LA}(\mathbf{x}) = \sum_{t=3}^T |x_t - \beta_1 x_{t-1} - \beta_2 x_{t-2}|$$

Denote the minimisers by $\hat{\beta}_{LS}$ and $\hat{\beta}_{LA}$ (calculated by `ARp.beta.est`), and define the estimated residuals to be $\hat{e}_t = x_t - \hat{\beta}_1 x_{t-1} - \hat{\beta}_2 x_{t-2}$ for $t = 3, \dots, T$, and let \bar{e} be the mean of these. The \hat{e}_t can be re-centered to have mean zero by defining $\hat{e}_t = \hat{e}_t - \bar{e}$. (Results for \hat{e}_t obtained by LS and LA can be calculated with `ARp.resid`).

1. Use the residual resampling bootstrap method to evaluate the relative performance of the two parameter estimators. Specifically, estimate the variance and bias of the two estimators.

You may use `ARp.filter` as a helper function in your resampling code. Use at least $B = 1500$ bootstrap samples, each as long as the original data sequence ($T = 100$). To do a resampling, initialise values for x_1 and x_2 by picking a random consecutive subsequence from the data.

The LS estimator is optimal for Gaussian AR(p) processes. Explain if it is also optimal for this problem?

2. Compute a 95% prediction interval for x_{101} based on both estimators. This means using the bootstrapped time series and parameter estimates obtained in part 1) to estimate the corresponding residual distribution and in turn use this to simulate a value x_{101} for the observed time series. Note that the variability of the simulated x_{101} values should reflect both our lack of knowledge about the parameter values and our lack of knowledge about the residual distribution. Then find the limits in the prediction interval for x_{101} as quantiles in the simulated x_{101} values.

Problem B: Permutation test

Bilirubin (see <http://en.wikipedia.org/wiki/Bilirubin>) is a breakdown product of haemoglobin, which is a principal component of red blood cells. If the liver has suffered degeneration, if the decomposition of haemoglobin is elevated, or if the gall bladder has been destroyed, large amounts of bilirubin can accumulate in the blood, leading to jaundice. The following data (taken from Jørgensen (1993)) contain measurements of the concentration of bilirubin (mg/dL) in blood samples taken from three young men.

Individual	Concentration (mg/dL)										
1	0.14	0.20	0.23	0.27	0.27	0.34	0.41	0.41	0.55	0.61	0.66
2	0.20	0.27	0.32	0.34	0.34	0.38	0.41	0.41	0.48	0.55	
3	0.32	0.41	0.41	0.55	0.55	0.62	0.71	0.91			

We will use the F-statistic to perform a permutation test.

Download the data file `bilirubin.txt` from the course webpage and read it into R using

```
bilirubin <- read.table("bilirubin.txt",header=T)

> head(bilirubin)
  meas pers
1 0.14  p1
2 0.20  p1
3 0.23  p1
4 0.27  p1
5 0.27  p1
6 0.34  p1
```

The first column, labelled `meas`, contains the concentrations (mg/dL) as shown in the table. The second column, `pers`, is an indicator for the individual.

1. Use a boxplot to inspect the logarithms of the concentrations for each individual. Be careful to use the same y-axis to make the plots comparable. Use the function `lm` in R to fit the regression model

$$\log Y_{ij} = \beta_i + \epsilon_{ij}, \quad \text{with } i = 1, 2, 3 \text{ and } j = 1, \dots, n_i \quad (1)$$

where $n_1 = 11$, $n_2 = 10$ and $n_3 = 8$, and $\epsilon_{ij} \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma^2)$. Use the F-test to test the hypothesis that $\beta_1 = \beta_2 = \beta_3$ and save the value of the F-statistic as `Fval`. Is the hypothesis accepted?

(Hint: Define the model as `m(log(meas) ~ pers, data=bilirubin)`. Then, the F-statistic of the test of interest is contained in the default output of `summary(lm)`)

2. Write a function `permTest()` which generates a permutation of the data between the three individuals, consequently fits the model given in (1) and finally returns the value of the F-statistic for testing $\beta_1 = \beta_2 = \beta_3$.
3. Perform a permutation test using the function `permTest` to generate a sample of size 999 for the F-statistic. Compute the p-value for `Fval` using this sample. What do you observe?

Problem C: The EM-algorithm and bootstrapping

Let x_1, \dots, x_n and y_1, \dots, y_n be independent random variables, where the x_i 's have an exponential distribution with intensity λ_0 and the y_i 's have an exponential distribution with intensity λ_1 . Assume we do not observe $x_1, \dots, x_n, y_1, \dots, y_n$ directly, but that we observe

$$z_i = \max(x_i, y_i) \quad \text{for } i = 1, \dots, n \quad (2)$$

and

$$u_i = I(x_i \geq y_i) \quad \text{for } i = 1, \dots, n, \quad (3)$$

where $I(A) = 1$ if A is true and 0 otherwise. Thus, for each $i = 1, \dots, n$ we observe the largest value of x_i and y_i and we know whether the observed value is x_i or y_i . Based on the observed $(z_i, u_i), i = 1, \dots, n$ we will use the EM algorithm to find the maximum likelihood estimates for (λ_0, λ_1)

1. Write down the log likelihood function for the complete data $(x_i, y_i), i = 1, \dots, n$. Use this to show that

$$\begin{aligned} E \left[\ln f(\mathbf{x}, \mathbf{y} | \lambda_0, \lambda_1) | \mathbf{z}, \mathbf{u}, \lambda_0^{(t)}, \lambda_1^{(t)} \right] &= n(\ln \lambda_0 + \ln \lambda_1) \\ &- \lambda_0 \sum_{i=1}^n \left[u_i z_i + (1 - u_i) \left(\frac{1}{\lambda_0^{(t)}} - \frac{z_i}{\exp\{\lambda_0^{(t)} z_i\} - 1} \right) \right] \\ &- \lambda_1 \sum_{i=1}^n \left[(1 - u_i) z_i + u_i \left(\frac{1}{\lambda_1^{(t)}} - \frac{z_i}{\exp\{\lambda_1^{(t)} z_i\} - 1} \right) \right] \end{aligned}$$

2. Using the EM algorithm, use the result you found in point 1) to find a recursion in $(\lambda_0^{(t)}, \lambda_1^{(t)})$ for finding the maximum likelihood estimates for (λ_0, λ_1) . Implement the recursion and find the maximum likelihood estimates when the data is as specified in tile files `z.txt` and `u.txt` available from the course home page. Visualise the convergence of the algorithm in a plot.
3. Use bootstrapping to estimate the standard deviations and the biases of each of $\hat{\lambda}_0$ and $\hat{\lambda}_1$ and to estimate $\text{Corr}[\hat{\lambda}_0, \hat{\lambda}_1]$. Present pseudocode for your bootstrap algorithm. Discuss briefly whether you would prefer the maximum likelihood estimates or the bias corrected estimates for λ_0 and λ_1 in this case.
4. For the situation defined here, you find an analytical formula for $f_{Z_i, U_i}(z_i, u_i | \lambda_0, \lambda_1)$?
Is it possible to find analytical formulas for the maximum likelihood estimators $\hat{\lambda}_0$ and $\hat{\lambda}_1$? Find the mle for $\hat{\lambda}_0$ and $\hat{\lambda}_1$ analytically or numerically. What are the advantages of optimizing the likelihood directly compared to the EM algorithm?

Literature

Jørgensen, B. (1993). The Theory of Linear Models. Chapman and Hall

Oral presentations

Date	Problem	Team
05.04	Ex3: Problem A	Group 6
05.04	Ex3: Problem B	Group 5
05.04	Ex3: Problem C	Group 7